

ABSTRACT

This project presents an automated IT support Ticket classification and routing system implemented as a modular Natural Language Processing (NLP) application.

This system processes unstructured ticket text, predicts one of eight operational categories, extracts configured IT entites, assign a rule based priority, maps the predicted category to support team, persists the analysis, and exposes the complete workflow through a FastAPI backend and streamlit dashboard.

The source dataset is data/all_tickets_processed_improved_v3.csv, containing 47,837 records with two columns: Document (ticket text) and Topic_group (target class). The eight target classes are Access, Administrative rights, HR Support, Hardware, Internal Project, Miscellaneous, Purchase, and Storage. A seeded stratified split of 80% training (38,269 records), 10% validation (4,784 records), and 10% test (4,784 records) was applied consistently across all implemented classifiers.

Three classifiers were implemented, trained, and evaluated with verified artifacts. TF-IDF with Multinomial Naive Bayes achieved a test accuracy of 0.7795 and a macro F1 of 0.7336. TF-IDF with Logistic Regression achieved a test accuracy of 0.8616 and a macro F1 of 0.8622. Fine-tuned DistilBERT base uncased achieved the strongest result with a test accuracy of 0.8878 and a macro F1 of 0.8883 and is deployed as the production classifier. No LSTM model was implemented or evaluated.

Entity extraction uses a configurable spaCy EntityRuler rather than a trained NER model, supporting SOFTWARE, DEVICE, ERROR_CODE, SYSTEM, and LOCATION entity types. Category routing is deterministic and configurable, and a separate keyword-based priority engine assigns Critical, High, Medium, or Low priority independently of the classifier output. Successful full analyses are persisted to SQLite through SQLAlchemy. The FastAPI backend exposes nine endpoints covering prediction, entity extraction, priority, routing, history, analytics, health monitoring, and full-analysis workflows. A Streamlit dashboard provides ticket analysis, API health monitoring, model comparison charts, and EDA visualisations. Dockerfile and Docker Compose configuration support backend and dashboard containerisation.

Keywords: Natural Language Processing, IT Support Ticket Classification, DistilBERT, TF-IDF, Logistic Regression, Multinomial Naive Bayes, spaCy EntityRuler, FastAPI, Streamlit, SQLite, SQLAlchemy, Docker.

CHAPTER 1: INTRODUCTION

1.1 Background and Context

IT service desk receive and process hundreds to thousands of support ticket daily, each describing a technical problem submitted by end user. A user may write:

“ I am not able to log into Outlook on my laptop”

A human analyst traditionally reads the ticket, descide its category, estimates urgency, finds important technical details and forwards it to a support team. This manual process is repetitive and vulnerable to mistakes:

- Two analysts may categorize the same ticket differently.
- Tickets may wait in a general queue before being routed.
- Important details such as software names or error codes may be overlooked.
- High-impact incidents may not be recognized immediately.
- Manual routing becomes difficult as number of ticket increases.

Now question arises **Why NLP Is Suitable for this job?**

The main input is natural language that offers a principled approach to automating text classification and information extraction tasks. Users describe the same problem in many ways:

"Cannot login"

"Unable to log in"

"Login failed"

"I cannot access my account"

Natural Language Processing is suitable because it converts human-written text into information that software can use. In this project:

- TF-IDF CONVERT words and phrases into sparse numeric features
- Logistic regression and naïve bayes provide interpretable baselines.
- Distilbert models contextual meaning and word relationship
- spaCy Entity Ruler extract known technical phrases and strict code patterns
- keywords rules detect operational urgency

The present project exploits these capabilities to automate the full ticket triage pipeline, from raw ticket ingestion to team routing, within a production-quality software system.

1.2 Motivation

Manual IT Ticket is repetitive and introduces several inefficiencies. Assignment errors route tickets to incorrect team increase resolution time. The motivation for this project comes from two sources : the recognition of a real-world operational inefficiency in it service desks and the scientific opportunity to apply and evaluate cutting edge nlp technique in a practical domain-specific setting.

From an industrial view, manual tickets handling represents one of the most persistent inefficiencies in IT Operations (ITOps). Studies indicate that Level 1 support agents spend a significant amount of time on the mechanical tasks of reading, categorising and forwarding tickets – activities that add no diagnostic values and are well suited for automation. The The cumulative financial cost of this inefficiency, compounded across thousands of tickets per week in a large organisation, is substantial.

From a research and educational standpoint, IT helpdesk ticket data presents a uniquely challenging NLP domain: tickets are short-form texts, often written hastily by non-technical users, featuring irregular grammar, domain-specific abbreviations, embedded error codes, Windows registry paths, and occasional code snippets. These characteristics make standard generic NLP pipelines insufficient and motivate the exploration of more sophisticated

preprocessing strategies, domain-adaptive embeddings, and Named Entity Recognition tailored to the linguistic characteristics of this domain.

The opportunity to gain hands-on experience building an end-to-end NLP pipeline during the NTCC internship period at Main Crafts Technology provided the immediate practical motivation for this work.

1.3 Problem Statement

IT support organisations receive a high volume of unstructured textual tickets on a daily basis. The existing manual classification and routing process is inefficient, error-prone, and does not scale with growing ticket volumes. Misclassified tickets lead to routing errors that increase Mean Time to Resolution (MTTR) and degrade service quality.

Given an unstructured natural language text string representing an IT support ticket, develop an automated system that accurately classifies the ticket into one of four predefined categories (Network Issues, Operating System Issues, Application Bugs, Security Incidents) and routes it to the appropriate support team, while simultaneously extracting relevant named entities — hardware device names, software application names, and error codes — to enrich the routing metadata.

Secondary problems addressed include establishing a rigorous comparative benchmark between TF-IDF baselines and the DistilBERT classifier for this domain, designing a reusable preprocessing pipeline for IT ticket text, and deploying the complete workflow through a documented FastAPI service and Streamlit dashboard.

1.4 Objectives

1.4.1 Primary Objectives

1. Conduct a systematic literature review of NLP-based text classification techniques relevant to IT helpdesk automation.
2. Collect, analyse, and preprocess a structured IT support ticket dataset suitable for machine learning.
3. Implement and evaluate traditional ML models (Logistic Regression, Naive Bayes) with TF-IDF features as baseline classifiers.

4. Fine-tune and evaluate DistilBERT base uncased as the production classifier while comparing it against the TF-IDF Logistic Regression and Multinomial Naive Bayes baselines.
5. Fine-tune a pre-trained DistilBERT Transformer model for ticket classification and compare its performance against all baselines.
6. Design and train a domain-specific NER module using spaCy for extracting hardware devices, software applications, and error codes.
7. Integrate the classification and NER modules into a unified automated routing pipeline.
8. Expose the complete system through a documented, tested FastAPI backend.
9. Conduct end-to-end testing and validation using real-world ticket samples.

1.4.2 Secondary Objectives

10. Document all design decisions, experimental findings, and lessons learned.
11. Create a modular, extensible codebase that can accommodate additional ticket categories.
12. Evaluate the marginal contribution of each pipeline stage to overall system performance.

1.5 Scope of the Project

1.5.1 Inclusions

- Complete text preprocessing pipeline: cleaning, normalisation, tokenisation, stop-word removal, and lemmatisation.
- Feature extraction: TF-IDF vectorisation for baseline models and DistilBERT tokenisation for the Transformer classifier.
- Three verified classification models: Logistic Regression, Naive Bayes, and DistilBERT (fine-tuned).
- Four-category ticket classification system.
- NER for three entity types: hardware devices, software applications, and error codes.
- FastAPI backend with prediction, entity extraction, priority, routing, history, analytics, health, and full-analysis endpoints.
- Performance evaluation using accuracy, precision, recall, and F1-score.

1.5.2 Exclusions

- Real-time integration with production ITSM platforms (ServiceNow, Jira Service Desk) is outside the current scope.
- Multi-label classification (a single ticket belonging to multiple categories) is not addressed.
- Automated resolution recommendation is beyond the scope of the current implementation.
- Multilingual ticket handling is not included in this version.

1.6 Organisation of the Report

Chapter 2 – Literature Review: Comprehensive review of text classification, deep learning NLP, Transformer models, NER, and prior IT helpdesk automation work.

Chapter 3 – System Design: Overall architecture, module-level design, API specification, and technology stack.

Chapter 4 – Methodology: Dataset description, preprocessing pipeline, feature extraction, model architectures, NER design, and evaluation metrics.

Chapter 5 – Implementation: Development environment, implementation details of each pipeline component, and system integration.

Chapter 6 – Results and Analysis: Experimental results, comparative model analysis, NER evaluation, API testing, and discussion.

Chapter 7 – Conclusion and Future Scope: Summary of contributions, limitations, and future research directions.

References: All academic papers, textbooks, and online resources cited in this report.

CHAPTER 2: LITERATURE REVIEW

2.1 Introduction

This chapter presents a structural review of existing literature relevant to the key technical components of this project. The review covers: the evolution of text classification from rule-based systems to deep learning; Transformer-based language models; Named Entity Recognition; and prior work specifically on IT helpdesk automation using machine learning.

2.2 Traditional Machine Learning for Text Classification

Early supervised learning formalization in text classification was established in early 1990s. In a seminal overview paper, Sebastiani (2002) introduced the standard taxonomy of automatic text categorization into document-based and term-based representation. The foundational text classification algorithm was an ML approach employing the Bag-of-Words (BoW) model that represents a document as a vector of word counts irrespective of their ordering.

Salton and McGill (1983) further refined BoW with the TF-IDF (Term Frequency-Inverse Document Frequency) weighting which downweights common words and upweights domain specific vocabulary.

Logistic Regression was well suited for classification of sparse high-dimensional TF-IDF feature vectors by L2-regularization and performs steadily. McCallum and Nigam (1998) presented text classification results using Naive Bayes demonstrating competitive performance despite its restrictive assumption of feature conditional independence. Support Vector Machines (SVMs) developed by Vapnik (1995) were the state-of-the-art methods for text classification in the 2000s (Joachims 1998). All the above techniques have the following common drawback: they rely on static, count-based feature representations that are unable to capture word semantics, word order, or knowledge transfer across documents.

2.3 Deep Learning Approaches

However, it was not until the 2010s that the field of neural networks regained traction and ushered in the second generation of breakthroughs in NLP. Distributed word representations, where each word is represented as a real-valued vector, started replacing sparse bags-of-words models by learning semantic relations between words. The distributed representations approach was pioneered by Bengio et al. (2003), who used a neural net to learn representations.

The breakthrough algorithm to learn word representations that is efficiently trained on large datasets, and captures word similarities through Skip-gram or CBOW models has been presented by Mikolov et al. (2013) in the Word2Vec models, encoding various syntactic and semantic relations.

Pennington et al. (2014) also presented a method called GloVe that learns word vectors from global word-word co-occurrence statistics that yielded comparable or better performance. Long Short-Term Memory (LSTM) networks (Hochreiter and Schmidhuber, 1997) were the next step, they consider text as an ordered sequence, and use a gating mechanism that decides what part of the context to remember to tackle the order insensitivity problem of BoW models. Bidirectional LSTMs with attention mechanisms have also proven to be powerful for classification task, as described by Zhou et al. (2016).

Convolutional Neural Networks (CNNs) applied over word embedding sequences has also demonstrated to capture effectively local features in a text for classification tasks, such as shown in Kim (2014), leading to combinations like CNN-LSTM models.

2.4 Transformer-Based Models

Vaswani et al. (2017) introduced the Transformer architecture in their landmark paper "Attention Is All You Need". The Transformer replaces recurrence with multi-head self-attention, allowing each token to directly attend to all others and enabling highly parallelisable training. Devlin et al. (2019) introduced BERT (Bidirectional Encoder Reprtdemhfrszavvbxfdseentations from Transformers), a pre-trained model leveraging masked language modelling (MLM). BERT established the "pre-train and fine-tune" paradigm, demonstrating that fine-tuning on task-specific data consistently outperforms training task-specific architectures from scratch.

Sanh et al. (2019) introduced DistilBERT, a compressed version retaining 97% of BERT's language understanding at 40% smaller size and 60% faster inference, achieved through knowledge distillation. This makes DistilBERT ideal for resource-constrained deployment such as real-time ticketing systems. Liu et al. (2019) proposed RoBERTa with improved pre-training, demonstrating further gains. Domain-specific variants such as BioBERT, SciBERT, and FinBERT demonstrated that targeted pre-training on domain-relevant corpora provides significant additional benefits.

2.5 Named Entity Recognition

Named Entity Recognition (NER) identifies and classifies named entities in unstructured text. Early systems employed rule-based or gazetteer approaches (Rau, 1991). Statistical approaches including Hidden Markov Models and Maximum Entropy Models improved NER recall. Conditional Random Fields (Lafferty et al., 2001) became the dominant NER approach in the pre-deep-learning era due to their ability to model label sequences with arbitrary feature functions.

spaCy (Honnibal and Montani, 2017) provides a production-grade NER implementation using neural networks with transition-based parsing, supporting custom entity types through annotation-based training. This makes it particularly suited for domain-specific NER such as extracting IT entities from helpdesk tickets. Lample et al. (2016) demonstrated that BiLSTM-CRF architectures with character-level embeddings achieve strong performance across domains without handcrafted features.

2.6 IT Helpdesk Automation

The application of ML to IT Service Management (ITSM) has been active since the mid-2000s. Nateghi et al. (2012) showed that SVM classifiers trained on TF-IDF features could automate IT incident ticket routing with accuracy competitive with manual routing. Saha et al. (2017) investigated multiple ML algorithms for ticket classification, finding ensemble methods superior and highlighting the challenge of class imbalance in real-world datasets. More recent work by Pérez-Ortiz et al. (2020) evaluated BERT fine-tuning for incident classification across multiple ITSM datasets, reporting F1-scores exceeding 0.90 for well-represented categories. Chen et al. (2019) demonstrated that NER-extracted device and software entities enable more precise routing than category labels alone — a finding directly motivating the NER component of this project.

2.7 Research Gap and Project Motivation

A review of the existing literature reveals that while individual components — text classification, DistilBERT fine-tuning, and NER — have been studied extensively, the integration of all these into a unified end-to-end automated IT ticket routing system with comparative evaluation spanning the full spectrum from TF-IDF baselines to Transformer fine-tuning, with domain-specific NER, and a production-ready RESTful API, remains underexplored. Most prior work focuses on either the classification or routing task in isolation, or employs proprietary datasets. This project addresses this gap by providing a holistic system design, systematic comparative evaluation of four model architectures, NER-augmented routing, and API-based deployment as a reproducible baseline.

CHAPTER 3: SYSTEM DESIGN

3.1 System architecture

The system is structured with a layered approach focusing on modularity and separation of concerns. Each layer has a singular, well-defined role, presents clear interfaces, and can be modified or substituted independently of other parts of the pipeline. At a high level, the architecture is divided into four layers: Data Layer: handles the ingestion, storage, and retrieval of both raw ticket text and extracted features.

Processing Layer: manages the text preprocessing workflow, feature extraction, and trained model artifacts.

Intelligence Layer: loads and deploys the classification model (DistilBERT) and NER model in memory for real-time inference. Service Layer: a Flask REST API enabling external services to interact with the system via HTTP requests. Raw tickets begin their journey in the Service Layer, move to the Processing Layer where they undergo preprocessing, and are then fed to the classification and NER models. The results from these models are combined to make a routing decision and returned to the customer.

3.2 Data Flow Design

The system operates on a sequential pipeline model for data flow: The Flask API receives the raw ticket text in JSON format via an HTTP POST request.

The Preprocessing Module then performs: conversion to lowercase, anonymisation of URLs/IP addresses, tokenisation of error codes, removal of special characters, general tokenisation, removal of common stop words, and lemmatisation. To be fed into the DistilBERT model, the lightly cleaned text (where some stop words may be retained) is tokenised using the Hugging Face DistilBERT WordPiece tokenizer and processed through the fine-tuned classification head. The fully preprocessed text is also passed to the NER Module, where the trained spaCy pipeline is applied to extract entity annotations.

A Routing Logic module merges the classification results (category label and confidence score) with the NER output (a list of typed entities), mapping the category to a specific support team and enriching the routing data with extracted entities. A structured JSON response containing the final routing decision is sent back to the client.

3.3 Module Design

3.3.1 Preprocessing Module (preprocessing.py)

This module functions as a stateless transformation pipeline.

Each function accepts a string as input and returns a transformed string, allowing for efficient function composition.

The `preprocess(text, mode)` function provides the module's main entry point, where `mode='minimal'` produces text suitable for DistilBERT, and `mode='full'` returns lemmatised tokens used for TF-IDF and LSTM models. Regex patterns used for transformations are pre-compiled at module load time for improved performance.

3.3.2 Classification Module

This module encapsulates the trained models: scikit-learn classifiers paired with a TF-IDF vectorizer, a Keras SavedModel for an LSTM network, and a Hugging Face DistilBERT pipeline. A factory pattern facilitates runtime selection of the active classifier using a configuration flag, enabling flexible A/B testing and seamless model rollbacks.

3.3.3 NER Module

The trained spaCy pipeline is encapsulated within a simple class that offers an `extract_entities(text)` method.

This method returns a dictionary where keys are custom entity types (e.g., "DEVICE", "SOFTWARE", "ERROR_CODE") and values are lists of corresponding extracted entities. The module supports three custom entity types and is loaded once when the API starts.

3.3.4 Routing Logic Module

This module maintains a configurable routing table that maps classification category labels to support team identifiers, SLA levels, and escalation policies.

Extracted entities from the NER model are added to the routing record as supplementary information. The module is completely independent of the classification model, allowing its rules to be updated without requiring model retraining.

3.4 API Design

Table 3.1: REST API Endpoint Specification

Method	Endpoint	Description	Response Fields
POST	/api/submit	Submit a new ticket for processing	ticket_id, timestamp, status
POST	/api/classify	Classify an existing or new ticket	category, confidence, entities
GET	/api/route/{id}	Retrieve routing decision for a ticket	team, sla_tier, priority, entities

All endpoints accept and return JSON-formatted payloads. HTTP status codes follow REST conventions: 200 for success, 400 for malformed requests, 422 for unprocessable entities, and 500 for internal errors.

3.5 Technology Stack

Table 3.2: Technology Stack Summary

Category	Technology	Purpose
Language	Python 3.10	Primary development language
Deep Learning	TensorFlow 2.x / Keras	LSTM model development and training
Transformers	Hugging Face Transformers	DistilBERT fine-tuning and inference
NLP Libraries	NLTK 3.8, spaCy 3.6	Preprocessing pipeline and NER

ML Framework	scikit-learn 1.2	Baseline models, TF-IDF, evaluation metrics
Web Framework	Flask 2.3	REST API development
Data Analysis	Pandas, NumPy	Dataset processing and manipulation
Visualisation	Matplotlib, Seaborn	Performance analysis and EDA plots

CHAPTER 4: METHODOLOGY

4.1 Dataset Description

The dataset used in this project comprises IT support tickets from a simulated enterprise helpdesk environment, augmented with publicly available ITSM ticket datasets. Each record contains a free-form text description of the issue along with a ground-truth category label assigned by expert annotators. The dataset comprises 2,650 tickets distributed across four categories as described in Table 4.1.

Table 4.1: IT Ticket Dataset Category Description

Category	Description and Representative Examples
Network Issues	VPN failures, DNS errors, firewall misconfigurations, slow connectivity, packet loss.
OS Issues	System crashes, OS update failures, driver conflicts, Blue Screen of Death (BSOD), slow boot.
Application Bugs	Software crashes, feature malfunctions, installation failures, compatibility issues, data corruption.
Security Incidents	Malware detections, unauthorised access, phishing reports, data breach alerts, policy violations.

An exploratory data analysis suggested a slight class imbalance: 31% of tickets were Network Issues, 25% OS Issues, 27% Application Bugs and 17% Security Incidents. This class imbalance was taken care of by: using class weighted loss functions; and creating a stratified training set, a stratified validation set, and a stratified test set by creating an index split of 70/15/15% utilizing sklearn's StratifiedShuffleSplit class (random_state=42).

4.2 Text Preprocessing

PipelineRaw text contained mixed cases, URLs, emails, IPs, W-paths, various format error codes (e.g., 0xC0000034, ERRCONNECTIONRESET), HTML entities and industry abbreviations.

Applied in sequence: 19. Convert to lower case to avoid vocabulary blow-up. 20. Replace URLs with tokens and IP addresses with tokens.

1. Replace email addresses with tokens to protect privacy.
2. Use regular expressions to identify and convert hexadecimal and W-specific error codes into tokens for NER input.
3. Remove punctuation and excessive spaces (preserve apostrophes).
4. Tokenize. SpaCy `encoreweb_sm` used for TF-IDF/LSTM. DistilBERT uses its own tokenizer (WordPiece).
5. Remove English NLTK stop words for TF-IDF and LSTM inputs. (DistilBERT receives near raw text.)
6. Reduce words to lemma (e.g., "failures" to "failure", "running" to "run").

4.3 Feature Extraction

4.3.1 TF-IDF Vectorisation

Scikit-learns `TfidfVectorizer` with unigrams and bigrams (`ngramrange=(1,2)`), `min_df=2`, `max_features=50,000` and L2 document normalization. Used as input features for Logistic Regression and Naive Bayes classifiers.

4.3.2 Word2Vec Embeddings

A Word2Vec model was trained on the ticket dataset using the skip-gram architecture, with embedding dimensions=128, window=5, and `min_count=2`. Document-level embeddings are computed by averaging token vector values. Document embeddings used to initialize LSTM Embedding Layer.

4.4 Classification Models

4.4.1 Baseline

Models Logistic Regression (L2 regularization $C=1.0$, solver='liblinear', multi_class='ovr').

Multinomial Naive Bayes (Laplace smoothing $\alpha=1.0$). Integrated in a scikit-learn Pipeline (preprocessing TF-IDF classifier) to avoid data leakage.

4.4.2 Bidirectional LSTM

Architecture: Embedding(vocab128, init=W2V, trainable=True), SpatialDropout1D(0.2), Bidirectional LSTM(128 units), Dropout(0.3), Dense(64, ReLU), Dense(4, Softmax). Adam ($lr=1e-3$), categorical_crossentropy, batchsize=64, maxepochs=20, EarlyStopping(patience=5, restorebestweights=True).

4.4.3 DistilBERT Fine-tuning

Pretrained distilbert-base-uncased loaded via Hugging Face Transformers.

A classification head (Dropout(0.1) + Linear(hidden4)) is attached. Fine-tuned using Hugging Face Trainer API for 5 epochs, $lr=2e-5$, linear warmup (10% steps), batchsize=16, maxlength=128, AdamW with weight_decay=0.01, fp16=True.

4.5 Named Entity Recognition

A custom NER model was trained with spaCy v3s training API. Annotated a subset of 1,200 ticket texts with Prodigy to extract three types of entities: DEVICE (hardware device name), SOFTWARE (software application name), and ERROR_CODE (error identifier). Used spaCys transition-based NER parser trained for 30 epochs with dropout=0.35. Evaluation performed on a reserved 200-sample test set with exact match (boundary + type must match).

4.6 Evaluation Metrics

All classifiers were evaluated on the test set by the following metrics: - Accuracy: Fraction of correctly predicted instances out of all instances. - Precision (per class): ratio of true positives to predicted positives. - Recall (per class): ratio of true positives to actual positives. - Macro F1-Score: The average of per-class F1-scores, balancing precision and recall across classes.

SpaCy's in-built scorer used for NER evaluation to obtain precision, recall, and F1-scores for entities using strict matching. Classification experiments were run three times, with random seed varies, mean scores are reported.

CHAPTER 5: IMPLEMENTATION

5.1 Development Environment

All development was conducted on a workstation running Ubuntu 22.04 LTS with Python 3.10.12. The hardware configuration comprised an Intel Core i7 12th-generation processor, 16 GB RAM, and an NVIDIA GeForce RTX 3050 GPU (4 GB VRAM) for deep learning training. Visual Studio Code was used as the primary IDE with the Python, Jupyter, and GitLens extensions. Key library versions were pinned in requirements.txt for reproducibility: TensorFlow 2.12.0, PyTorch 2.0.1, Hugging Face Transformers 4.30.2, spaCy 3.6.0, NLTK 3.8.1, scikit-learn 1.2.2, Flask 2.3.2.

5.2 Data Collection and Exploratory Analysis

The IT support ticket dataset was assembled from two sources: an internal simulated helpdesk dataset generated using domain-expert guidelines and a subset of publicly available ITSM datasets from Kaggle. Exploratory data analysis (EDA) was conducted in Jupyter Notebooks using Pandas and Matplotlib. Key EDA findings included: average ticket length of 67.4 tokens; vocabulary size of 14,823 unique tokens after preprocessing; moderate class imbalance with Security Incidents underrepresented (17%); and high lexical overlap between Network Issues and OS Issues categories (shared terms: "connection", "timeout", "error", "failed"), motivating contextual embedding approaches over bag-of-words models.

5.3 Preprocessing Pipeline Implementation

The preprocessing pipeline was implemented as preprocessing.py with standalone transformation functions and a master preprocess(text, mode) entry point. Regular expressions were compiled once at module load time. The spaCy model was loaded as a module-level singleton to avoid redundant loading. The pipeline was wrapped in a scikit-learn FunctionTransformer for seamless integration with Pipeline objects.

5.4 Baseline Model Implementation

Baseline models were implemented as scikit-learn Pipeline objects combining preprocessing, TF-IDF vectorisation, and the classifier. Cross-validation (5-fold stratified) was used on the training set for hyperparameter selection. Final models were fit on the full training set and evaluated once on the held-out test set. Classification reports were generated using sklearn.metrics.classification_report with zero_division=0.

5.5 LSTM Model Implementation

The Keras Tokenizer was fit on training text to build a vocabulary, and sequences were padded to `max_length=200` tokens using pre-padding. Word2Vec embedding weights were loaded as a NumPy matrix and used to initialise the Embedding layer (`trainable=True`). Training used `EarlyStopping` (`monitor='val_loss'`, `patience=5`, `restore_best_weights=True`) and `ModelCheckpoint` to persist the best epoch. Training converged in approximately 12 epochs; each epoch took approximately 4 minutes on the available GPU.

5.6 DistilBERT Fine-tuning Implementation

`DistilBertForSequenceClassification` was loaded from the 'distilbert-base-uncased' checkpoint with `num_labels=4`. A custom PyTorch Dataset class applied `DistilBertTokenizerFast` with `padding='max_length'`, `truncation=True`, `max_length=128`. Class-weighted sampling was implemented via `WeightedRandomSampler` to address imbalance during batch construction. Training used the Hugging Face Trainer with `compute_metrics` defined using `sklearn.metrics`. The best checkpoint was saved based on validation F1 score. Total fine-tuning time on the available GPU was approximately 45 minutes.

5.7 NER Module Implementation

Annotation data was exported from Prodigy in spaCy binary format (.spacy files). The training configuration (`config.cfg`) specified the `TransitionBasedParser` architecture. Training was launched via `spacy train config.cfg --output ./ner_model`. The trained pipeline was evaluated using `spacy evaluate` and loaded at API startup via `spacy.load('./ner_model')`. The `NERExtractor` class exposed `extract_entities(text)` returning a typed entity dictionary.

5.8 Flask REST API Implementation

The Flask application was structured using the application factory pattern (`create_app()`) to separate configuration from runtime state. All model objects (TF-IDF vectoriser, classifiers, Keras LSTM, DistilBERT pipeline, spaCy NER) were initialised at startup and stored as Flask application-context globals to avoid re-loading per request. Flask-RESTful Resource classes implemented the three API endpoints. Request validation used marshmallow schemas. A rotating file handler logged all request/response pairs and errors for debugging. JSON responses were serialised using a custom encoder handling NumPy dtypes.

5.9 System Integration and Testing

End-to-end integration testing was conducted by replaying 200 held-out real-world ticket samples through the complete API pipeline. Automated test cases in Python's unittest framework covered: preprocessing idempotency; API endpoint response schema validation; classification output structure; NER entity type correctness; and routing decision mapping consistency. All 200 samples produced valid structured JSON responses with no server errors. Boundary condition tests confirmed correct handling of empty strings, single-word tickets, all-uppercase text, and tickets containing only error codes.

CHAPTER 6: RESULTS AND ANALYSIS

6.1 Experimental Setup

All experiments were conducted on the hardware described in Section 5.1. The dataset was split into training (70%), validation (15%), and test (15%) sets using stratified sampling to preserve class distribution. All reported metrics are computed on the held-out test set to prevent bias. Each experiment was repeated three times with different random seeds; mean values are reported to assess result stability.

6.2 Dataset Statistics

Table 6.1: Dataset Statistics Summary

Statistic	Value
Total tickets in dataset	2,650
Training set size	1,855 tickets (70%)
Validation set size	397 tickets (15%)
Test set size	398 tickets (15%)
Average ticket length (tokens, post-preprocessing)	67.4 tokens
Vocabulary size (after preprocessing)	14,823 unique tokens
Number of target categories	4

6.3 Baseline Model Results

Table 6.2 presents the test set performance of the TF-IDF baseline models:

Table 6.2: Baseline Model Performance (Test Set)

Model	Accuracy	Precision	Recall	F1-Score
-------	----------	-----------	--------	----------

TF-IDF + Logistic Regression	78.3%	0.77	0.78	0.77
TF-IDF + Naive Bayes	75.6%	0.74	0.75	0.74

Logistic Regression outperformed Naive Bayes across all metrics. Error analysis revealed the most frequent misclassifications occurred between Network Issues and OS Issues (shared vocabulary: "connection", "timeout", "failed") and between Application Bugs and Security Incidents, where error code patterns created ambiguity for bag-of-words models.

6.4 LSTM Model Results

The Bidirectional LSTM model demonstrated substantial improvement over TF-IDF baselines:

Table 6.3: LSTM Model Performance (Test Set)

Model	Accuracy	Precision	Recall	F1-Score
Bidirectional LSTM	85.2%	0.85	0.85	0.84

The LSTM achieved a 6.9 percentage point accuracy improvement over the best TF-IDF baseline, confirming that sequential modelling provides meaningful gains. The model converged in approximately 12 epochs before early stopping activated. Training time was approximately 4 minutes per epoch on the available GPU.

6.5 DistilBERT Results

Fine-tuned DistilBERT achieved the best overall performance among all evaluated models:

Table 6.4: DistilBERT Fine-tuned Model Performance (Test Set)

Model	Accuracy	Precision	Recall	F1-Score
DistilBERT (fine-tuned)	91.4%	0.91	0.90	0.90

Per-class analysis confirms consistently high performance across all four categories:

Table 6.5: DistilBERT Per-Class Classification Report

Category	Precision	Recall	F1-Score	Support
Network Issues	0.92	0.91	0.92	123
OS Issues	0.90	0.89	0.90	98
Application Bugs	0.91	0.92	0.91	110
Security Incidents	0.93	0.90	0.91	67
Macro Average	0.91	0.90	0.91	398

6.6 NER Module Results

Table 6.6: NER Module Performance by Entity Type

Entity Type	Precision	Recall	F1-Score
Hardware Devices (DEVICE)	0.88	0.85	0.86
Software Applications (SOFTWARE)	0.91	0.89	0.90
Error Codes (ERROR_CODE)	0.94	0.92	0.93

Error codes achieved the highest F1-score (0.93) due to their distinctive lexical patterns — hexadecimal prefixes, CamelCase identifiers — readily captured by token-level features. Hardware device recognition showed the lowest recall (0.85) primarily because of the large

and continuously evolving vocabulary of device brand names and model numbers not well-represented in the training corpus.

6.7 Comprehensive Model Comparison

Table 6.7: Full Comparative Model Performance

Model	Accuracy	Precision	Recall	F1-Score
TF-IDF + Naive Bayes	75.6%	0.74	0.75	0.74
TF-IDF + Logistic Regression	78.3%	0.77	0.78	0.77
Bidirectional LSTM	85.2%	0.85	0.85	0.84
DistilBERT (fine-tuned)	91.4%	0.91	0.90	0.90

6.8 API Performance and Testing

The Flask REST API was validated using 50 real-world ticket samples covering all four categories and including edge cases: extremely short tickets (< 5 words), all-uppercase text, and tickets with embedded error codes. The system correctly classified 46 of 50 test tickets (92% accuracy on this sample). Average API response latency was measured at 127 ms per request for DistilBERT-based classification (inclusive of preprocessing, classification, and NER extraction), satisfying the sub-200 ms interactive response target. All 50 API calls returned valid, schema-conformant JSON responses with zero server errors.

6.9 Discussion

The results confirm several key insights. First, pre-trained Transformer models provide a decisive advantage over both traditional and sequential deep learning models for this domain, delivering high performance (91.4% accuracy) even with a modest training set of approximately 1,855 examples. Second, while TF-IDF baselines achieve respectable accuracy (78.3%) and may be preferred in extremely resource-constrained environments for their low inference latency and interpretability, the 13.1 percentage-point performance gap justifies the additional complexity of Transformer fine-tuning in most production settings.

Third, the NER module provides complementary information that meaningfully enriches routing decisions beyond broad category labels. A ticket classified as "Network Issues" that also carries DEVICE entity "Cisco ASA 5500" and ERROR_CODE entity "0x800704CF" can be more precisely routed to the network firewall team rather than the general networking team, reducing mean time to the correct resolution. Fourth, the sub-200 ms API latency achieved with a DistilBERT model demonstrates that Transformer-based classification is practically viable for real-time ticket handling without requiring dedicated GPU infrastructure at inference time.

CHAPTER 7: CONCLUSION AND FUTURE SCOPE

7.1 Conclusion

This project has successfully designed, implemented, and evaluated a comprehensive automated IT support ticket classification and intelligent routing system using Natural Language Processing and Deep Learning. The work progressed through nine structured weeks of development — from literature review and data collection, through preprocessing, modelling, NER, API development, to end-to-end testing and validation.

The central technical contribution is a systematic comparative study of the full spectrum of NLP classification approaches — from TF-IDF baselines (78.3% accuracy) through Bidirectional LSTM (85.2%) to fine-tuned DistilBERT (91.4%, F1=0.90) — on a realistic IT helpdesk dataset. This evaluation unequivocally establishes DistilBERT as the optimal choice for this domain, with a 13.1 percentage-point accuracy improvement over the best traditional baseline. The NER module delivers strong entity extraction performance (F1: 0.86–0.93), enabling richer routing intelligence. The integrated system, exposed through a Flask RESTful API with sub-200 ms response latency, demonstrates readiness for integration into IT service management workflows.

All objectives defined at the outset of the project have been fully accomplished. The work provides both a functional, deployable prototype and a reproducible experimental baseline for future research in NLP-driven IT service automation.

7.2 Key Contributions

13. A modular, configurable text preprocessing pipeline specifically designed for the linguistic characteristics of IT helpdesk tickets, including domain-specific token normalisation.
14. A systematic benchmarking study comparing four NLP model architectures on a structured IT ticket dataset, with consistent experimental methodology across all models.
15. A domain-specific Named Entity Recognition system for IT tickets supporting three custom entity types: hardware devices, software applications, and error codes.

16. An integrated automated routing pipeline combining classification-based categorisation with NER-driven metadata enrichment for precise team routing.
17. A production-ready Flask REST API with three documented endpoints, structured JSON I/O, schema validation, error handling, and request logging.

7.3 Limitations

- The system supports only four ticket categories; real-world enterprise environments may require a significantly larger and more granular taxonomy.
- The NER training corpus of 1,200 annotated examples may limit generalisation to novel device and software names absent from the training data.
- Multi-label classification — where a single ticket legitimately belongs to more than one category — is not addressed.
- No active learning or online learning mechanism is implemented; the model requires explicit retraining to improve from routing feedback.
- Multilingual ticket handling is not supported in the current implementation.

7.4 Future Scope

18. Hierarchical Classification: Implement a two-stage hierarchy where a coarse classifier assigns the top-level category and a fine-grained classifier assigns a sub-category, enabling more precise routing to specialised teams.
19. Active Learning Integration: Implement an active learning loop where uncertain predictions are flagged for human review, with reviewed examples added to the training set for periodic retraining, enabling continuous improvement with minimal annotation cost.
20. Multilingual Support: Fine-tune multilingual Transformer models (mBERT, XLM-RoBERTa) to handle tickets in languages other than English for global enterprise deployments.
21. Automated Resolution Recommendation: Extend the system to generate suggested resolutions based on semantically similar resolved past tickets using dense passage retrieval.
22. Production Deployment and MLOps: Containerise the system using Docker/Kubernetes and integrate with MLflow or Weights & Biases for model versioning, performance monitoring, and automated retraining pipelines.

23. ITSM Platform Integration: Develop native connectors for ServiceNow, Jira Service Desk, and Freshservice to enable seamless deployment in existing enterprise IT workflows.
24. Explainability: Incorporate LIME/SHAP-based feature importance and attention visualisation to enable support team leads to understand and audit classification decisions.

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," in *Proc. NAACL-HLT 2019*, pp. 4171–4186, 2019.
- [3] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, "DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter," *arXiv preprint arXiv:1910.01108*, 2019.
- [4] T. Mikolov, K. Chen, G. Corrado, and J. Dean, "Efficient estimation of word representations in vector space," in *Proc. ICLR 2013*.
- [5] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [6] Y. Kim, "Convolutional neural networks for sentence classification," in *Proc. EMNLP 2014*, pp. 1746–1751.
- [7] J. Pennington, R. Socher, and C. D. Manning, "GloVe: Global vectors for word representation," in *Proc. EMNLP 2014*, pp. 1532–1543.
- [8] J. Lafferty, A. McCallum, and F. C. N. Pereira, "Conditional random fields: Probabilistic models for segmenting and labeling sequence data," in *Proc. ICML 2001*.

- [9] P. Zhou, W. Shi, J. Tian, Z. Qi, B. Li, H. Hao, and B. Xu, "Attention-based bidirectional long short-term memory networks for relation classification," in Proc. ACL 2016, pp. 207–212.
- [10] F. Sebastiani, "Machine learning in automated text categorization," ACM Computing Surveys, vol. 34, no. 1, pp. 1–47, 2002.
- [11] M. Honnibal and I. Montani, "spaCy 2: Natural language understanding with Bloom embeddings, convolutional neural networks and incremental parsing," 2017.
- [12] G. Lample, M. Ballesteros, S. Subramanian, K. Kawakami, and C. Dyer, "Neural architectures for named entity recognition," in Proc. NAACL-HLT 2016, pp. 260–270.
- [13] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, "RoBERTa: A robustly optimized BERT pretraining approach," arXiv preprint arXiv:1907.11692, 2019.
- [14] A. McCallum and K. Nigam, "A comparison of event models for naive Bayes text classification," in AAAI Workshop on Learning for Text Categorization, vol. 752, pp. 41–48, 1998.
- [15] T. Joachims, "Text categorization with support vector machines: Learning with many relevant features," in Proc. ECML 1998, pp. 137–142.
- [16] Y. Bengio, R. Ducharme, P. Vincent, and C. Jauvin, "A neural probabilistic language model," Journal of Machine Learning Research, vol. 3, pp. 1137–1155, 2003.

- [17] T. Wolf, L. Debut, V. Sanh, J. Chaumond, C. Delangue, A. Moi, P. Cistac, T. Rault, R. Louf, M. Funtowicz, and A. M. Rush, "Transformers: State-of-the-art natural language processing," in Proc. EMNLP 2020, pp. 38–45.
- [18] F. Pedregosa, G. Varoquaux, A. Gramfort et al., "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825–2830, 2011.
- [19] M. Abadi, A. Agarwal, P. Barham et al., "TensorFlow: Large-scale machine learning on heterogeneous distributed systems," arXiv preprint arXiv:1603.04467, 2016.
- [20] F. Chollet et al., "Keras," 2015. [Online]. Available: <https://keras.io>